

Course Code, CSE, IIT Jammu

MTech I (1st semester)

Design and Analysis of Algorithms

Lab Assignment No 1 - Introduction and Advanced Data Structures

AUTUMN Semester 2019-20

Instructions:

1. The date of submission **will be specified only two days before**. Therefore, right from the time the assignment is uploaded the students must start implementing the assignments.
 2. For every delayed submission beyond the deadline, 10 marks per day will be deducted from the maximum marks of the assignment, without any exception, whatsoever may be the scapegoat.
 3. You can use any programming language, except Python.
 4. For every program whether specified or not, it is necessary to time your program on input dataset of a large size and give a critique of the theoretical asymptotic analysis of the algorithm that your program is based on and the empirical timing analysis of the program that you have written. The a-priori estimates and the a-posteriori analysis must be in sync with each other.
 5. All the assignments must be submitted in the form of a zip file containing the Program Source, the screenshot of the output that you obtained, the DataSet/Input Test data used and a ReadMe .txt file explaining what platform to use, what are the input parameters required for execution and how to execute it. Also write your conclusion in ReadMe file.
 6. Perform usual error checking. Dont go overboard on this, but dont let your program die because of divide by zero.
 7. Remember, your programs could be checked by a code-cheating program, so please follow the code of academic integrity.
 8. The viva for this assignment will be taken on a future date as specified.
 9. **Maximum Points 200.**
-

1. Write a program for Insertion sort, Merge sort, Quick sort and Heap sort **to count number of exchanges/swaps required** to sort the input data. Consider the worst case, the best case and the average case inputs to test your program. Use appropriate dataset that contains atleast 10000 values to test your program.
2. Modify the program in Assignment#1 to measure time required to sort the input data. Discuss what could be various approaches to **time** the programs, including the utilities available in Linux. Use an approach that gives you the finest resolution in time.
3. Now, assume that the new requirement is that (if not already taken care of) - your program must output the **time in microseconds**. How to achieve the same ? Whatever approach you are able to devise as the solution to this question, compare your solution with the program **microresolutiontimer.c** uploaded on the LMS. Give a critique of your comparison.
4. Write a program to implement Binary Exponent Algorithm explained in class, time the program and analyze as explained in the instructions. Compare your implementation with the *Square-and-multiply* method of exponentiation.
5. This assignment concerns the Completely Fair Scheduling algorithm used in Linux that employs Red Black Trees as the underlying data structure. The Completely Fair Scheduler (CFS) is a process scheduler that handles the CPU resource allocation for executing processes, and aims to maximize overall CPU utilization while also maximizing the interactive performance. The CFS scheduler implementation is not based as in traditionally schedulers on **run queues**. Instead, a **redblack tree** implements a "timeline" of future task executions. Additionally, the scheduler uses nanosecond granularity accounting, the atomic units by which an individual process' share of the CPU was allocated (thus making redundant the previous notion of timeslices). This precise knowledge also means that no specific heuristics are required to determine the interactivity of a process, for example. [Source: Wiki.]

For this assignment, read the accompanying notes on CFS uploaded on the LMS. Implement the CFS scheduler to illustrate one of the applications of the Red Black Trees.

6. Write a program that illustrates the comparison of an ordinary Binary tree with any one of the 2-3, AVL, or Red Black trees and illustrates the specific benefits of one of the 2-3, AVL, or Red Black trees that you have used in your implementation. Obviously for the implementation and fairer comparison, use the same data set as input. Then, extend this program to compare the 2-3, AVL, or Red Black trees with each other. Analyze the comparison and give concrete conclusions.
7. For this program, you are to fill in the functions, writing versions of (1) insertion sort (2) merge sort and (3) myownsort. myownsort() is a variant of the existing sorting algorithms that is to be devised by you with an attempt to improve upon the performance of the merge sort. Use the skeleton of the codes including here for writing your programs - filling in the functions required. The sorting functions that you write must be modeled on the illustration of the Bubble sort shown. Assume N is the input array that you use for evaluating the programs.

After successfully writing and exeuting the code, answer the following questions:

- (a) For each algorithm, how large (approximately) can N be if the sort must take no more than 10 seconds?
- (b) Can you think of a quick way to improve bubble sort by a factor of 2? Write a "semibubbleSort(a)" function to do so.

Your programs must obey the following template:

```
// Program 7 — some sorting
// compile: gcc -o program7 program7.c
// usage:    program7 100
//           (replace "100" by the array size you want)
//           Name:
//
// Date submitted:
//
/* include the required files */
// Declarations
void bubbleSort(int a[]);
void insertSort(int a[]);
void mergeSort(int a[]);
void mySort(int a[]);
void fillRandom(int a[]);
void check(int a[]);

int N = 0;                                // Global for array size

int main(int argc, char *argv[]) {
    int a[1000000];                        // array to sort, N < a million
    timeval t;                             // for timing the sorts
    int starttime, endtime; // for timing the sorts

    N = atoi(argv[1]);                     // set size of array

    // Time a bubble sort
    fillRandom(a);
    gettimeofday(&t, NULL);
    starttime = t.tv_sec;
    bubbleSort(a);
    gettimeofday(&t, NULL);
    endtime = t.tv_sec;
    printf("Bubble Sort time = %f%d%d", (endtime - starttime));
    check(a);

    // time an insertion sort
    fillRandom(a);
    gettimeofday(&t, NULL);
    starttime = t.tv_sec;
    insertSort(a);
    gettimeofday(&t, NULL);
```

```

        endtime = t.tv_sec;
        printf("Insertion Sort time = %f%d%d", (endtime - starttime));
        check(a);

        // time a merge sort
        fillRandom(a);
        gettimeofday(&t, NULL);
        starttime = t.tv_sec;
        mergeSort(a);
        gettimeofday(&t, NULL);
        endtime = t.tv_sec;
        printf("Merge Sort time = %f%d%d", (endtime - starttime));
        check(a);

        // time a "my" sort
        fillRandom(a);
        gettimeofday(&t, NULL);
        starttime = t.tv_sec;
        mySort(a);
        gettimeofday(&t, NULL);
        endtime = t.tv_sec;
        printf("MySort time = %f%d%d", (endtime - starttime));
        check(a);

        return 0;
    }

    void bubbleSort(int a[]) {
        for(int i = 0; i <= N; i++)
            for(int j = 0; j < N-1; j++)
                if(a[j] > a[j+1]){
                    int tmp = a[j];
                    a[j] = a[j+1];
                    a[j+1] = tmp;
                }
        return;
    }

    void insertSort(int a[]) {
    }

    void mergeSort(int a[]) {
    }

    void mySort(int a[]) {
    }

    void fillRandom(int a[]) {
        for(int i = 0; i < N; i++)
            a[i] = rand();
        return;
    }

    void check(int a[]) {
        for(int i = 0; i < N-1; i++)
            if(a[i] > a[i+1]){
                printf("Failed to sort. See item: %d \n\n", i) ;
                return;
            }
        printf ("Check passed \n\n");
    }

```

}

8. Write and fill-in the statements in the sketch of the program shown below, that allows a user to interact with a heap. The heap will never have more than 20 elements in it. You should write your own additional functions to help perform these heap operations. When the heap is in "heapstate = 0," then it is just a heap, and no attempt should be made to maintain either a min or max heap property thereafter, unless a "min" or "max" command is issued. For consistency, implement the algorithms in the fashion given in the text CLRS, so that the heap will look the same for everybody after each command.

Sample Transcript: add 1

Added 1

1-

add 5

Added 5

1-5-

add 2

Added 2

1-5-2-

sort

Sorting

1-5-2-

max

Making a max heap

5-1-2-

sort

Sorting

1-2-5-

min

Making a min heap

1-2-5-

sort

Sorting

5-2-1-

```
// Program 8 — some heap practice
// compile: gcc -o program8 program8.c
// usage:   ./program8
// Name:
//
// Date submitted:
//
```

```
#include required files
```

```
#define N 20
```

```
// Global Variables
```

```
int heapstate = 0; // 0 = heap, 1 = maxheap, 2 = minheap
```

```
void add(int A[], int heapsize, int val){
    printf("Added %d\n", val) ;
}
```

```
int extract(int A[], int &heapsize){
    printf("Extracting .....%n");
}
```

```

}

void make_max_heap(int A[], int heapsize){
    printf("Making a max heap %n");
}

void make_min_heap(int A[], int heapsize){
    printf("Making a min heap %n");
}

void heapsort(int A[], int heapsize){
    if(heapstate == 0) return;
    printf(" %n");
}

void print(int A[], int heapsize){
    for(int i = 1; i <= heapsize; i++)
        printf("%d -%n", A[i]);
}

int main(){
    int A[N+1];
    int heapsize = 0;
    char *command = "";
    int val = 0;

    // Initialize
    // Remember that our arrays start with index 1.
    for(int i = 1; i <= N; i++)
        A[i] = 0;

    do{
        scanf("%s", command);
        if(command == "add"){
            // Add a value to the end of the heap.
            // Should maintain the heap property according to heapstate.
            scanf("%d", &val);
            add(A, heapsize, val);
        }

        if(command == "extract"){
            // Should extract A[1] and maintain heap property,
            // according to heapstate
            printf("Extracted %- %n", extract(A, heapsize));
        }

        if(command == "neither"){
            heapstate = 0;
        }

        if(command == "max"){
            heapstate = 1;
            make_max_heap(A, heapsize);
        }

        if(command == "min"){
            heapstate = 2;
            make_min_heap(A, heapsize);
        }

        if(command == "sort"){
            // Sort into increasing order if heapstate = 1,
            // and into decreasing order if heapstate = 2.

```

```

        heapsort(A, heapsize);
    }

    print(A, heapsize);
} while (command != "end");
}

```

9. You are given an undirected graph, and you should return the number of connected components in the graph. The graph will be given as a list of vertices, followed by a list of edges, as shown in the sample transcript below. Each list is terminated with the word "end."

Your output should give the number of components, and then a sorted list of sizes of the components, from largest to smallest. Each vertex name will be a single character; in fact, a lower-case letter. Please make sure that your output matches the sample transcript shown below, for the ease of the grading. You do not need to do error checking on the input.

Sample Transcript:c

\$ program9

Please input your vertices:

a
b

10. We have looked at the theoretical analysis of algorithms using asymptotic notations. In practice, however, a theoretically inferior algorithm (I) can work better than a theoretically superior algorithm (S) due to the constants hidden in the big-Oh analysis. Also, several optimizations that can be done for a simpler algorithm may not quite be applicable for S. However, in all cases, S should eventually outperform I. Finally, whether inferior or superior, blackboard algorithms do not consider all possible situations, something which a programmer has to deal with.

So, what is the Task at hand in this assignment ?

Well, in this programming assignment, you will solve the problem of multiplying two polynomials (not necessarily of the same degree) using two different methods.

- (a) The first method, call it simple, is the straightforward quadratic method.
- (b) The second method, call it $n\log 3$, is a divide-and-conquer (D&C) method which runs in $\theta(n\log 3)$. You will subdivide the input polygons based on even or odd indices rather than the first half and second half method. Clearly you must first figure out the algorithm

Input Format Please read this section carefully and make sure that you understand it clearly. You will read the input from three input files. The first two files contain two sets of polynomials to be multiplied. You may assume that they are called *first.in* and *second.in*.

Each of the two files contains a set of coefficients that describes a set of polynomials as follows.

- The coefficients are listed in ascending order of power.
- Each line represents exactly one real number coefficient written in ASCII.
- The number of coefficients is not expected to be small. The precision is expected to be up to three decimal points. For example, if first.in is in the following format, then this means that it describes the set of polynomials 15 , $15 + 0.5x$, and $15 + 0.5x + x^2$.

15
0.5
1

- The third input file will list a set of indices of coefficients (starting from zero for the polynomial of degree zero) that needs to be output each on a separate line. You may assume that the coefficient\’s file is called coeff.in. For example, the following file lists the indices 0, 2, and 1.

0
2
1

Your task is best illustrated by an example as follows. Assume first.in and coeff.in are as before, and second.in is as follows.

1
1

That is second.in describes the set of polynomials 1 and $1 + x$.

Your task will be to read the input files one line from each of the three files at a time, compose the two polynomials up to that line, multiply the two polynomials using the two methods, output the running time of the two methods (in milliseconds), and output the value of the corresponding coefficient which is listed on the same line of coeff.in (the exact output format will be described in Output Format Section).

For example, for the previous input files you will do the following.

- (a) Read the polynomial 15 of degree 0 from first.in, the polynomial 1 of degree 0 from second.in, and the coefficient index 0 from coeff.in. Multiply the two polynomials 15 and 1 using the two methods. Output the running time for the two methods. Output the value of the coefficient index 0 of the resultant polynomial which in this case equals to 15
- (b) Read the polynomial $15+0.5x$ of degree 1 from first.in, the polynomial $1+x$ of degree 1 from second.in, and the coefficient index 2 from coeff.in. Multiply the two polynomials $15+0.5x$ and $1+x$ using the two methods. Output the running time for the two methods. Output the value of the coefficient index 2 of the resultant polynomial which in this case equals to 0.5.
- (c) Read the polynomial $15+0.5x+x^2$ from first.in, the polynomial $1+x+0x^2$ from second.in, and the coefficient index 1 from coeff.in. Notice that second.in has already reached the end of file marker in the previous step, so you have to handle the case of polynomials of different lengths as well as polynomials of equal lengths. Also, notice that coeff.in will always have number of lines equal to the number of lines in the longer file of first.in and second.in (the file for the longer polynomial). Multiply the two polynomials $15+0.5x+x^2$ and $1+x+0x^2$ using the two methods. Output the running time for the two methods. Output the value of the coefficient index 1 of the resultant polynomial which in this case equals to 15.5.
- (d) Now first.in, second.in, and coeff.in have reached the end of file marker, and your task is now complete.

Output Format Please read this section carefully since any deviation in the output format will result in loss of points. Your output should consist of two parts as we discussed in previous section. The corresponding values of the coefficients listed in coeff.in, AND the running time of the two methods, both as text files. You may assume that the output files are called coeff.out and running.out respectively. The running.out file should consist of three columns text file. The first column represents n , which is the degree of the multiplied polynomials (the longer one), the running time for simple in milliseconds, and the running time for $n\log 3$ in milliseconds. This file is intended to be used to draw a plot in order to show the asymptotic performance of both algorithms for each value of n and then see if it actually matches the theoretical results.

For example, running.out can be as follows.

```
0 1000 800
1 2000 1700
.
.
.
100 50000 43700
```

The coeff.out file should be the actual answers to the input given above. Specifically, It should list the values of the required coefficients based on the coeff.in input file as we described in Section 4. You should output each of the coefficients as a real number on a separate line with a precision up to three decimal points. For example, the output for the example we discussed in previous section should be in the following format.

```
15
0.5
15.5
```